

# EDAN95

## Applied Machine Learning

### Lecture 7: Sequence Prediction

Pierre Nugues

Pierre.Nugues@cs.lth.se  
[http://cs.lth.se/pierre\\_nugues/](http://cs.lth.se/pierre_nugues/)

November 23, 2020

# Outline

In the previous lecture, we used networks to produce one output  $y$  per input vector  $\mathbf{x}$ , for instance one category per sentence.

Given an input sequence  $\mathbf{x}$ , we will now produce an output sequence:  $\mathbf{y}$ .

We will experiment three kinds of neural networks:

- 1 Feed forward
- 2 Recurrent
- 3 LSTM

In the laboratory assignment, you will use the two last ones.

# Motivation

The analysis of sentences often involves the analysis of words.

We can divide it in three main tasks:

- 1 Identify the type of word, for instance noun or verb using the classical grammar;
- 2 Identify a group or segment, for instance are these three words, *Kjell Olof Andersson*, the name of a person;
- 3 Identify the relations between two words: for instance is this group the subject of a verb? This corresponds to parsing, semantic analysis, or information extraction.

We will consider the two first tasks.

This lecture will show you how to solve the first one, part-of-speech tagging, and you will write a program for the second one, named entity recognition (NER), in the next laboratory assignment.

# Word Categorization: The Parts of Speech

Sentence:

*That round table might collapse*

Annotation:

| Words           | Parts of speech | POS tags |
|-----------------|-----------------|----------|
| <b>that</b>     | Determiner      | DET      |
| <b>round</b>    | Adjective       | ADJ      |
| <b>table</b>    | Noun            | NOUN     |
| <b>might</b>    | Modal verb      | AUX      |
| <b>collapse</b> | Verb            | VERB     |

The automatic annotation uses predefined POS tagsets such as the Penn Treebank tagset for English

# Ambiguity

| Words           | Possible tags  | Example of use  |
|-----------------|--|---|
| <b>that</b>     | Subordinating conjunction<br>Determiner<br>Adverb<br>Pronoun<br>Relative pronoun | <i>That he can swim is good</i><br><i>That white table</i><br><i>It is not that easy</i><br><i>That is the table</i><br><i>The table that collapsed</i> |
| <b>round</b>    | Verb<br>Preposition<br>Noun<br>Adjective<br>Adverb                               | <i>Round up the usual suspects</i><br><i>Turn round the corner</i><br><i>A big round</i><br><i>A round box</i><br><i>He went round</i>                  |
| <b>table</b>    | Noun<br>Verb   | <i>That white table</i><br><i>I table that</i>  |
| <b>might</b>    | Noun<br>Modal verb   | <i>The might of the wind</i><br><i>She might come</i>   |
| <b>collapse</b> | Noun<br>Verb   | <i>The collapse of the empire</i><br><i>The empire can collapse</i>   |

# Training Sets: The CoNLL Format

The CoNLL format is a tabular format to distribute annotated texts. This format was created for evaluations carried out by the Conference in natural language learning

The CoNLL annotation has varied much across the years. We use CoNLL-U, the latest iteration.

Annotation of the Spanish sentence:

*La reestructuración de los otros bancos checos se está acompañando por la reducción del personal*

*'The restructuring of Czech banks is accompanied by the reduction of personnel'*

# Example of Annotation (CoNLL-U)

*La reestructuración de los otros bancos checos se está acompañando por la reducción del personal*

| ID | FORM             | LEMMA            | UPOS  | FEATS  |
|----|------------------|------------------|-------|--|
| 1  | La               | el               | DET   | Definite=Def Gender=Fem Number=Sing PronType=Art       |
| 2  | reestructuración | reestructuración | NOUN  | Gender=Fem Number=Sing                                 |
| 3  | de               | de               | ADP   | AdpType=Prep   |
| 4  | los              | el               | DET   | Definite=Def Gender=Masc Number=Plur PronType=Art      |
| 5  | otros            | otro             | DET   | Gender=Masc Number=Plur PronType=Ind                   |
| 6  | bancos           | banco            | NOUN  | Gender=Masc Number=Plur                                |
| 7  | checos           | checo            | ADJ   | Gender=Masc Number=Plur                                |
| 8  | se               | se               | PRON  | Case=Acc Person=3 PrepCase=Npr PronType=Prs Reflex=Yes |
| 9  | está             | estar            | AUX   | Mood=Ind Number=Sing Person=3 Tense=Pres VerbForm=Fin  |
| 10 | acompañando      | acompañar        | VERB  | VerbForm=Ger   |
| 11 | por              | por              | ADP   | AdpType=Prep   |
| 12 | la               | el               | DET   | Definite=Def Gender=Fem Number=Sing PronType=Art       |
| 13 | reducción        | reducción        | NOUN  | Gender=Fem Number=Sing                                 |
| 14 | del              | del              | ADP   | AdpType=Preppron                                       |
| 15 | personal         | personal         | NOUN  | Gender=Masc Number=Sing                                |
| 16 | .                | .                | PUNCT | PunctType=Peri   |

# Another Example

| ID | FORM       | LEMMA      | PLEMMA     | POS   | PPOS  | FEAT | PFEAT |
|----|------------|------------|------------|-------|-------|------|-------|
| 1  | Battle     | battle     | battle     | NN    | NN    | —    | —     |
| 2  | -          | -          | -          | HYPH  | HYPH  | —    | —     |
| 3  | tested     | tested     | tested     | NN    | NN    | —    | —     |
| 4  | Japanese   | japanese   | japanese   | JJ    | JJ    | —    | —     |
| 5  | industrial | industrial | industrial | JJ    | JJ    | —    | —     |
| 6  | managers   | manager    | manager    | NNS   | NNS   | —    | —     |
| 7  | here       | here       | here       | RB    | RB    | —    | —     |
| 8  | always     | always     | always     | RB    | RB    | —    | —     |
| 9  | buck       | buck       | buck       | VBP   | VB    | —    | —     |
| 10 | up         | up         | up         | RP    | RP    | —    | —     |
| 11 | nervous    | nervous    | nervous    | JJ    | JJ    | —    | —     |
| 12 | newcomers  | newcomer   | newcomer   | NNS   | NNS   | —    | —     |
| 13 | with       | with       | with       | IN    | IN    | —    | —     |
| 14 | the        | the        | the        | DT    | DT    | —    | —     |
| 15 | tale       | tale       | tale       | NN    | NN    | —    | —     |
| 16 | of         | of         | of         | IN    | IN    | —    | —     |
| 17 | the        | the        | the        | DT    | DT    | —    | —     |
| 18 | first      | first      | first      | JJ    | JJ    | —    | —     |
| 19 | of         | of         | of         | IN    | IN    | —    | —     |
| 20 | their      | their      | their      | PRP\$ | PRP\$ | —    | —     |
| 21 | countrymen | countryman | countryman | NNS   | NNS   | —    | —     |
| 22 | to         | to         | to         | TO    | TO    | —    | —     |
| 23 | visit      | visit      | visit      | VB    | VB    | —    | —     |
| 24 | Mexico     | mexico     | mexico     | NNP   | NNP   | —    | —     |
| 25 | ,          | ,          | ,          | ,     | ,     | —    | —     |
| 26 | a          | a          | a          | DT    | DT    | —    | —     |
| 27 | boatload   | boatload   | boatload   | NN    | NN    | —    | —     |
| 28 | of         | of         | of         | IN    | IN    | —    | —     |
| 29 | samurai    | samurai    | samurai    | NN    | NN    | —    | —     |
| 30 | warriors   | warrior    | warrior    | NNS   | NNS   | —    | —     |
| 31 | blown      | blow       | blow       | VRN   | VRN   | —    | —     |



# Designing a Part-of-Speech Tagger

We will now create part-of-speech taggers, where we will examine three architectures:

- ➊ A feed-forward pipeline with a one-hot encoding of the words;
- ➋ A feed-forward pipeline with word embeddings: We will replace the one-hot vectors with GloVe embeddings;
- ➌ A recurrent neural network, either a simple RNN or a LSTM, with word embeddings.

# Features for Part-of-Speech Tagging

The word *visit* is ambiguous in English:

*I paid a **visit** to a friend → noun*

*I went to **visit** a friend → verb*

The context of the word enables us to tell, here an article or the infinitive marker

To train and apply the model, the tagger extracts a set of features from the surrounding words, for example, a sliding window spanning five words and centered on the current word.

We then associate the feature vector  $(w_{i-2}, w_{i-1}, w_i, w_{i+1}, w_{i+2})$  with the part-of-speech tag  $t_i$  at index  $i$ .

# Part-of-Speech Tagging

| ID  | FORM       | PPOS  |                |
|-----|------------|-------|----------------|
|     | BOS        | BOS   | Padding        |
|     | BOS        | BOS   |                |
| 1   | Battle     | NN    |                |
| 2   | -          | HYPH  |                |
| 3   | tested     | NN    |                |
| ... | ...        | ...   |                |
| 17  | the        | DT    |                |
| 18  | first      | JJ    |                |
| 19  | of         | IN    |                |
| 20  | their      | PRP\$ |                |
| 21  | countrymen | NNS   | Input features |
| 22  | to         | TO    |                |
| 23  | visit      | VB    | Predicted tag  |
| 24  | Mexico     |       | ↓              |
| 25  | ,          |       |                |
| 26  | a          |       |                |
| 27  | boatload   |       |                |
| ... | ...        | ...   |                |
| 34  | years      |       |                |
| 35  | ago        |       |                |
| 36  | .          |       |                |
|     | EOS        |       | Padding        |
|     | EOS        |       |                |

# Feature Vectors

| ID  | Feature vectors |            |            |            |            |           |           | PPOS  |
|-----|-----------------|------------|------------|------------|------------|-----------|-----------|-------|
|     | $w_{i-2}$       | $w_{i-1}$  | $w_i$      | $w_{i+1}$  | $w_{i+2}$  | $t_{i-2}$ | $t_{i-1}$ |       |
| 1   | BOS             | BOS        | Battle     | -          | tested     | BOS       | BOS       | NN    |
| 2   | BOS             | Battle     | -          | tested     | Japanese   | BOS       | NN        | HYPH  |
| 3   | Battle          | -          | tested     | Japanese   | industrial | NN        | HYPH      | JJ    |
| ... | ...             | ...        | ...        | ...        | ...        | ...       | ...       | ...   |
| 19  | the             | first      | of         | their      | countrymen | DT        | JJ        | IN    |
| 20  | first           | of         | their      | countrymen | to         | JJ        | IN        | PRP\$ |
| 21  | of              | their      | countrymen | to         | visit      | IN        | PRP\$     | NNS   |
| 22  | their           | countrymen | to         | visit      | Mexico     | PRP\$     | NNS       | TO    |
| 23  | countrymen      | to         | visit      | Mexico     | ,          | NNS       | TO        | VB    |
| 24  | to              | visit      | Mexico     | ,          | a          | TO        | VB        | NNP   |
| 25  | visit           | Mexico     | ,          | a          | boatload   | VB        | NNP       | ,     |
| ... | ...             | ...        | ...        | ...        | ...        | ...       | ...       | ...   |
| 34  | ashore          | 375        | years      | ago        | .          | RB        | CD        | NNS   |
| 35  | 375             | years      | ago        | .          | EOS        | CD        | NNS       | RB    |
| 36  | years           | ago        | .          | EOS        | EOS        | NNS       | RB        | .     |

# Architecture 1: A Feed-Forward Neural Network

We first use a feed-forward architecture corresponding to a logistic regression:

```
np.random.seed(0)

model = models.Sequential([Dense(NB_CLASSES,
                                   input_dim=X.shape[1],
                                   activation='softmax')])

model.compile(loss='sparse_categorical_crossentropy',
              optimizer=OPTIMIZER,
              metrics=['accuracy'])

model.summary()

model.fit(X, y, epochs=EPOCHS, batch_size=BATCH_SIZE)

model.save('out.model')
```

# Encoding the $\mathbf{y}$ Vector

In the previous examples, we used `categorical_crossentropy`. This requires that all the targets are encoded with one-hot vectors. For instance:

- determiner: [1, 0, 0, 0]
- noun: [0, 1, 0, 0]
- verb: [0, 0, 1, 0]
- adjective: [0, 0, 0, 1]

With `sparse_categorical_crossentropy`, we can use numerical indices:

- determiner: 1
- noun: 2
- verb: 3
- adjective: 4

We do not need to use the `to_categorical` function.

# Preprocessing

Preprocessing is more complex though: Four steps:

- 1 Read the corpus

```
train_sentences, dev_sentences, test_sentences, \
    column_names = load_ud_en_ewt()
```

- 2 Store the rows of the CoNLL corpus in dictionaries

```
conll_dict = CoNLDDictorizer(column_names, col_sep='\t')
train_dict = conll_dict.transform(train_sentences)
test_dict = conll_dict.transform(test_sentences)
```

- 3 Extract the features and store them in dictionaries

```
context_dictorizer = ContextDictorizer()
context_dictorizer.fit(train_dict)
X_dict, y_cat = context_dictorizer.transform(train_dict)
```

- 4 Vectorize the symbols

```
# We transform the X symbols into numbers
dict_vectorizer = DictVectorizer()
X_num = dict_vectorizer.fit_transform(X_dict)
```

# Code Example

Jupyter Notebook: [4.1-nn-pos-tagger.ipynb](#)



# Architecture 2: Using Embeddings

We replace the one-hot vectors with embeddings, the rest being the same  
Word embeddings are dense vectors obtained by a principal component analysis or another method.

They can be trained by the neural network or pretrained  
In this implementation:

- ① We use pretrained embeddings from the GloVe project;
- ② Our version of GloVe is lowercased, so we set all the characters in lowercase;
- ③ We add the embeddings as an `Embedding` layer at the start of the network;
- ④ We initialize the embedding layer with GloVe and make it trainable or not.

It would be possible to use a randomly initialized matrix as embeddings instead

# The Embedding Layer

```
model = models.Sequential([
    Embedding(cnt_uniq, EMBEDDING_DIM,
              input_length=2 * W_SIZE + 1),
    Flatten(),
    Dense(NB_CLASSES, activation='softmax')
])

if embedding_matrix is not None:
    model.layers[0].set_weights([embedding_matrix])
model.layers[0].trainable = True

model.compile(loss='sparse_categorical_crossentropy',
              optimizer=OPTIMIZER,
              metrics=['accuracy'])


model.summary()

model.fit(X, y, epochs=EPOCHS, batch_size=BATCH_SIZE)
```

# Code Example

Jupyter Notebook: `4.2-nn-pos-tagger-embeddings.ipynb`

# The RNN Architecture



# Input Format for RNNs

The input format is different from feed forward networks.

We need to build two lists: one for the input and the other for the output

|          |     |        |         |     |      |
|----------|-----|--------|---------|-----|------|
| <b>y</b> | DET | NOUN   | VERB    | DET | NOUN |
| <b>x</b> | The | waiter | brought | the | meal |

All the vectors in a same batch must have the same length. We pad them:

|          |     |     |     |     |        |         |     |      |
|----------|-----|-----|-----|-----|--------|---------|-----|------|
| <b>y</b> | PAD | PAD | PAD | DET | NOUN   | VERB    | DET | NOUN |
| <b>x</b> | PAD | PAD | PAD | The | waiter | brought | the | meal |

We could apply the padding after

# Building the Sequences

```
def build_sequences(corpus_dict, key_x='form', key_y='pos',
                    tolower=True):
    X, Y = [], []
    for sentence in corpus_dict:
        x, y = [], []
        for word in sentence:
            x += [word[key_x]]
            y += [word[key_y]]
        if tolower:
            x = list(map(str.lower, x))
        X += [x]
        Y += [y]
    return X, Y
```

At this point, we have **x** and **y** vectors of symbols

# Building Index Sequences

0 is for the padding symbol and 1 for the unknown words

```
idx_word = dict(enumerate(vocabulary_words, start=2))  
idx_pos = dict(enumerate(pos, start=2))  
word_idx = {v: k for k, v in idx_word.items()}  
pos_idx = {v: k for k, v in idx_pos.items()}
```

At this point, we have **x** and **y** vectors of numbers

# Padding the Index Sequences

We build the complete **X\_idx** and **Y\_idx** matrices for the whole corpus  
And we pad the matrices:

```
X = pad_sequences(X_idx)
```

```
Y = pad_sequences(Y_idx)
```

```
# The number of POS classes and 0 (padding symbol)
```

```
Y_train = to_categorical(Y, num_classes=len(pos) + 2)
```

`pad_sequences` can have an argument that specifies the maximal length  
`maxlen (MAX_SEQUENCE_LENGTH)`.

The padded sentences must have the same length in a batch. This is  
automatically computed by Keras



# Recurrent Neural Networks (RNN)

```
model = models.Sequential([
    Embedding(len(vocabulary_words) + 2,
              EMBEDDING_DIM,
              mask_zero=True,
              input_length=None),
    SimpleRNN(100, return_sequences=True),
    # Bidirectional(SimpleRNN(100, return_sequences=True)),
    Dense(NB_CLASSES + 2, activation='softmax')])

model.layers[0].set_weights([embedding_matrix])
# The default is True
model.layers[0].trainable = True
```

# Parameters

Keras functions have many parameters.

In case of doubt, read the documentation

A few useful parameters:

- 1 `mask_zero=True` is to tell whether or not the input value 0 is a special “padding” value;
- 2 `return_sequences=True` tells whether to return the last output in the output sequence, or the full sequence. In sequences, it is essential;
- 3 `recurrent_dropout=0.3` tells how much to drop for the linear transformation of the recurrent state.

# Code Example

Jupyter Notebook: `4.3-rnn-pos-tagger.ipynb`

# Long Short-Term Memory (LSTM)

```
model = models.Sequential([
    Embedding(len(vocabulary_words) + 2,
              EMBEDDING_DIM,
              mask_zero=True,
              input_length=None),
    Bidirectional(LSTM(100, return_sequences=True)),
    Dense(NB_CLASSES + 2, activation='softmax')])

model.layers[0].set_weights([embedding_matrix])
# The default is True
model.layers[0].trainable = True
```

# Segment Recognition

## Group detection – chunking –:

Brackets: [<sub>NG</sub> The government <sub>NG</sub>] has [<sub>NG</sub> other agencies and instruments <sub>NG</sub>] for pursuing [<sub>NG</sub> these other objectives <sub>NG</sub>] .

Tags: *The/I government/I has/O other/I agencies/I and/I instruments/I for/O pursuing/O these/I other/I objectives/I ./O*

Brackets: Even [<sub>NG</sub> Mao Tse-tung <sub>NG</sub>] [<sub>NG</sub> 's China <sub>NG</sub>] began in [<sub>NG</sub> 1949 <sub>NG</sub>] with [<sub>NG</sub> a partnership <sub>NG</sub>] between [<sub>NG</sub> the communists <sub>NG</sub>] and [<sub>NG</sub> a number <sub>NG</sub>] of [<sub>NG</sub> smaller, non-communists parties <sub>NG</sub>] .

Tags: *Even/O Mao/I Tse-tung/I 's/B China/I began/O in/O 1949/I with/O a/I partnership/I between/O the/I communists/I and/O a/I number/I of/O smaller/I ,/I non-communists/I parties/I ./O*

# Segment Categorization

Tages extendible to any type of chunks: nominal, verbal, etc.

For the IOB scheme, this means tags such as I.Type, O.Type, and B.Type, Types being NG, VG, PG, etc.

In CoNLL 2000, ten types of chunks

| Word           | POS | Group | Word             | POS | Group |
|----------------|-----|-------|------------------|-----|-------|
| <i>He</i>      | PRP | B-NP  | <i>to</i>        | TO  | B-PP  |
| <i>reckons</i> | VBZ | B-VP  | <i>only</i>      | RB  | B-NP  |
| <i>the</i>     | DT  | B-NP  | <i>£</i>         | #   | I-NP  |
| <i>current</i> | JJ  | I-NP  | <i>1.8</i>       | CD  | I-NP  |
| <i>account</i> | NN  | I-NP  | <i>billion</i>   | CD  | I-NP  |
| <i>deficit</i> | NN  | I-NP  | <i>in</i>        | IN  | B-PP  |
| <i>will</i>    | MD  | B-VP  | <i>September</i> | NNP | B-NP  |
| <i>narrow</i>  | VB  | I-VP  | <i>.</i>         | .   | O     |

Noun groups (NP) are in red and verb groups (VP) are in blue.

# IOB Annotation for Named Entities

| CoNLL 2002 |                | CoNLL 2003 |     |        |                |
|------------|----------------|------------|-----|--------|----------------|
| Words      | Named entities | Words      | POS | Groups | Named entities |
| Wolff      | B-PER          | U.N.       | NNP | I-NP   | I-ORG          |
| ,          | O              | official   | NN  | I-NP   | O              |
| currently  | O              | Ekeus      | NNP | I-NP   | I-PER          |
| a          | O              | heads      | VBZ | I-VP   | O              |
| journalist | O              | for        | IN  | I-PP   | O              |
| in         | O              | Baghdad    | NNP | I-NP   | I-LOC          |
| Argentina  | B-LOC          | .          | .   | O      | O              |
| ,          | O              |            |     |        |                |
| played     | O              |            |     |        |                |
| with       | O              |            |     |        |                |
| Del        | B-PER          |            |     |        |                |
| Bosque     | I-PER          |            |     |        |                |
| in         | O              |            |     |        |                |
| the        | O              |            |     |        |                |
| final      | O              |            |     |        |                |
| years      | O              |            |     |        |                |
| of         | O              |            |     |        |                |
| the        | O              |            |     |        |                |
| seventies  | O              |            |     |        |                |
| in         | O              |            |     |        |                |
| Real       | B-ORG          |            |     |        |                |
| Madrid     | I-ORG          |            |     |        |                |
| .          | O              |            |     |        |                |



# Named Entities: Proper Nouns

William Shakespeare

was born and brought

up in


Stratford-upon-Avon





# Others Entities: Common Nouns

Meeting with our guest on the landing at  
lunchtime



# Evaluation

There are different kinds of measures to evaluate the performance of machine learning techniques, for instance:

- Precision and recall in information retrieval and natural language processing;
- The *receiver operating characteristic* (ROC) in medicine.

|                   | Positive examples: $P$ | Negative examples: $N$ |
|-------------------|------------------------|------------------------|
| Classified as $P$ | True positives: $A$    | False positives: $B$   |
| Classified as $N$ | False negatives: $C$   | True negatives: $D$    |

More on the receiver operating characteristic here: [http://en.wikipedia.org/wiki/Receiver\\_operating\\_characteristic](http://en.wikipedia.org/wiki/Receiver_operating_characteristic)

# Recall, Precision, and the F-Measure

The **accuracy** is  $\frac{|A \cup D|}{|P \cup N|}$ .

**Recall** measures how much relevant examples the system has classified correctly, for  $P$ :

$$\text{Recall} = \frac{|A|}{|A \cup C|}.$$

**Precision** is the accuracy of what has been returned, for  $P$ :

$$\text{Precision} = \frac{|A|}{|A \cup B|}.$$

















Recall and precision are combined into the **F-measure**, which is defined as the harmonic mean of both numbers:

$$F = \frac{2 \cdot \text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}}.$$

# Evaluation

Accuracy, precision, and recall.

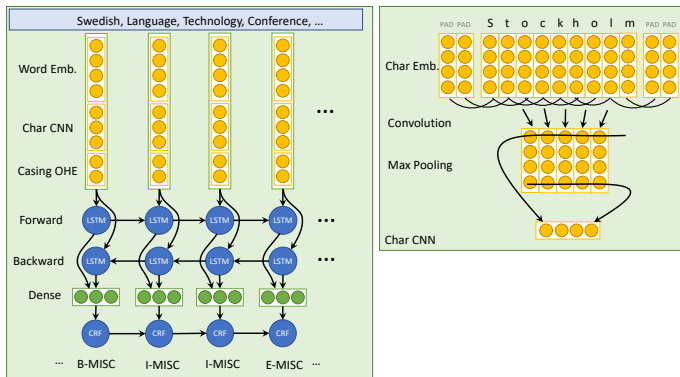
For noun groups with the predicted output:

| Word   | POS | Group       | Word   | POS | Group       |
|--|-----|-------------|--|-----|-------------|
|  <i>He</i>      | PRP | B-NP        |  <i>to</i>        | TO  | B-PP        |
|  <i>reckons</i> | VBZ | B-VP        |  <i>only</i>      | RB  | B-NP        |
|  <i>the</i>     | DT  | B-NP        |  <i>£</i>         | #   | I-NP        |
|  <i>current</i> | JJ  | <b>B-NP</b> |  <i>1.8</i>       | CD  | <b>B-NP</b> |
|  <i>account</i> | NN  | I-NP        |  <i>billion</i>   | CD  | I-NP        |
|  <i>deficit</i> | NN  | I-NP        |  <i>in</i>        | IN  | B-PP        |
|  <i>will</i>    | MD  | B-VP        |  <i>September</i> | NNP | B-NP        |
|  <i>narrow</i>  | VB  | I-VP        |  <i>.</i>         | .   | O           |

Accuracy =  $\frac{14}{16}$ , recall =  $\frac{2}{4} = 0.5$ , precision =  $\frac{2}{6} = 0.33$

harmonic mean =  $2 \times \frac{0.33 \times 0.5}{0.33 + 0.5} = 0.4$

# The Architecture of a Full-Fledged Network



Courtesy: Marcus Kiang. See also:

- Xuezhe Ma, Eduard Hovy, End-to-end Sequence Labeling via Bi-directional LSTM-CNNs-CRF, 2016, <https://arxiv.org/abs/1603.01354>
- Jason P.C. Chiu, Eric Nichols, Named Entity Recognition with Bidirectional LSTM-CNNs, 2016, <https://arxiv.org/abs/1511.08308>
- Guillaume Lample, Miguel Ballesteros, Sandeep Subramanian, Kazuya Kawakami, Chris Dyer, Neural Architectures for Named Entity Recognition, 2016, <https://arxiv.org/abs/1603.01360>